

Kolla / Habanero / Ubuntu 16.04

Introduction

This document explains how to deploy a Newton OpenStack instance on Power 8.

Configuration

Start from a Power system running Ubuntu 16.04 with two NICs . On our system, enP1p3s0f0 is the 'public' interface and enP1p3s0f0 the 'management' interface.

```
vi /etc/network/interfaces

auto enP1p3s0f0
iface enP1p3s0f0 inet static
    dns-nameservers 8.8.8.8
    address 172.31.250.2/24
    mtu 1500

auto enP1p9s0
iface enP1p9s0 inet static
    gateway 192.168.154.1
    dns-nameservers 8.8.8.8
    address 192.168.154.2/24
    mtu 1500
```

Preparation

Note: all the Kolla installation must be done under 'root'. Modify the /etc/ssh/sshd_config file.

```
ssh ubuntu@192.168.154.2

sudo vi /etc/ssh/sshd_config
# PermitRootLogin prohibit-password
PermitRootLogin yes

# PermitEmptyPasswords no

# PasswordAuthentication no

sudo service ssh restart

sudo cp /home/ubuntu/.ssh/authorized_keys /root/.ssh/

exit
```

```
ssh root@192.168.154.2
```

Update /etc/hosts

```
192.168.154.2 hnode1-11.pub.pic2.ibm.com hnode1-11
```

Add SMT off to /etc/rc.local

```
vi /etc/rc.local  
    ppc64_cpu --smt=off
```

Update

```
apt-get update  
apt-get upgrade -y  
apt-get install -y bridge-utils
```

Setup the 'public' interface as a bridge

```
vi /etc/network/interfaces  
  
# auto enP1p3s0f0  
# iface enP1p3s0f0 inet static  
#     address 172.31.250.2/24  
#     dns-nameservers 8.8.8.8  
#     mtu 1500  
  
auto enP1p3s0f0  
iface enP1p3s0f0 inet manual  
  
auto br1  
iface br1 inet static  
    address 172.31.250.2/24  
    gateway 172.31.255.254  
    dns-nameservers 8.8.8.8  
    bridge_ports enP1p3s0f0  
    bridge_fd 9  
    bridge_hello 2  
    bridge_maxage 12  
    bridge_stp off  
    mtu 1500
```

Remove gateway from enP1p9s0

```
auto enP1p9s0  
iface enP1p9s0 inet static  
    # gateway 192.168.154.1
```

```
dns-nameservers 8.8.8.8
address 192.168.154.2/24
mtu 1500
```

Reboot

```
reboot
```

Install pip

```
apt install python-pip -y
pip install --upgrade pip
```

Install Docker

```
apt-get install docker.io=1.10.3-0ubuntu6 -y
```

Enable insecure access to the registry

```
vi /etc/default/docker
DOCKER_OPTS="--insecure-registry 192.168.154.2:4000"
```

Restart Docker

```
service docker stop && service docker start
```

Update package

```
pip install -U docker-py
pip install -U python-openstackclient
pip install -U python-neutronclient
apt-get install -y python-dev libffi-dev libssl-dev gcc git
```

Create the drop-in unit directory for docker.service

```
mkdir -p /etc/systemd/system/docker.service.d
```

Create the drop-in unit file

```
tee /etc/systemd/system/docker.service.d/kolla.conf << 'EOF'
[Service]
MountFlags=shared
EOF
```

Reload and restart Docker

```
systemctl daemon-reload
systemctl restart docker
```

Mount

```
mount --make-shared /run
```

Install NTP

```
apt-get install ntp -y
```

Install Ansible

```
pip install -U ansible==2.1.2.0
```

Install Kolla

As Kolla on Power is a work in progress, the current version is a temporary custom version. It's distributed as a self contained images:

```
kolla-11182016.tar
kolla-etc-11182016.tar
kolla-tools-11182016.tar
kolla-usr-11182016.tar
```

The `kolla-11182016.tar` file will expand in your local home directory under `kolla`. It contains the original Kolla code from github.

```
tar -xvf kolla-11182016.tar
pip install kolla/
```

Make sure that Kolla version is '2.0.0.0rc2.dev624'

```
pip list | fgrep kolla
kolla (2.0.0.0rc2.dev624)
```

The `kolla-etc-11182016.tar` will expand in the `etc` local directory. It contains the Kolla deployment files specific to your environment. After configuration, these files need to be moved under the `/etc/kolla` directory.

```
tar -xvf kolla-etc-11182016.tar
mkdir -p /etc/kolla
cp etc/kolla/* /etc/kolla
```

The `kolla-tools-11182016.tar` will expand in your local directory. It's a set of tools to ease the build, deployment, cleanup tasks. They have to be reconfigured to match your local settings.

```
tar -xvf kolla-tools-11182016.tar
```

Note: after a system reboot run the `after-reboot` script to setup the environment. For automatic setup you can add the content of this file to your `/etc/rc.local` file.

Note: most of the tools don't have the IP addresses properly configured. Before running them, please make sure that the IP addresses match your environment.

The `kolla-usr-11182016.tar` will expand to the `usr` local directory. This is the modified Kolla code. It need to be moved under `/usr`

```
tar -xvf kolla-usr-11182016.tar
mkdir -p /usr/local/share/kolla
cp -R usr/local/share/kolla/* /usr/local/share/kolla
```

Create an Ubuntu base image for POWER

```
apt-get install -y debootstrap

curl -o debootstrap.sh
https://raw.githubusercontent.com/docker/docker/master/contrib/m
kimage/debootstrap

chmod 755 debootstrap.sh

./debootstrap.sh ubuntu --components=main,universe xenial

tar -C ubuntu -c . | sudo docker import - ubuntu:16.04

docker tag ubuntu:16.04 ubuntu:ppc64

docker images

docker tag <image id> 192.168.154.2:4000/ubuntu:ppc64

I.e docker tag b7ed97810eff 192.168.154.2:4000/ubuntu:ppc64
```

Install and start registry

```
apt-get install docker-registry -y

vi /etc/docker/registry/config.yml
http:
  addr: :4000
  headers:
    X-Content-Type-Options: [nosniff]

nohup sudo /usr/bin/docker-registry
/etc/docker/registry/config.yml &
```

This will start the registry service on port 4000. Check that it's opened for request

```
nc -zv <registry_host>:<port>

i.e.

nc -zv 192.168.154.2 4000
```

```
Connection to 192.168.154.2 4000 port [tcp/*] succeeded!
```

Push the Ubuntu ppc64 image in the local registry

```
docker push 192.168.154.2:4000/ubuntu:ppc64
```

Horizon fix

```
vi
/usr/local/share/kolla/ansible/roles/horizon/templates/local_settings
.j2
```

```
AVAILABLE_THEMES = [
    ('default', 'Default', 'themes/default'),
]
```

Disable libvirt profile

```
# Ubuntu
service libvirt-bin stop
update-rc.d libvirt-bin disable

sudo apparmor_parser -R /etc/apparmor.d/usr.sbin.libvirtd
```

Disable apparmor

```
service apparmor stop
update-rc.d -f apparmor remove
```

Open port 22

```
iptables -A INPUT -p tcp --dport ssh -j ACCEPT
```

VGA

In the image meta data, make sure to have vga set for the hw_video_model property

```
glance image-update cb5b0f93-5068-4ba8-b1b4-7b3673178d2b
--property hw_video_model=vga
```

Make sure to set the console on the kernel command line in the image itself. Make that change in /etc/default/grub, then rebuild the grub.cfg

```
GRUB_CMDLINE_LINUX="crashkernel=auto console=tty0"
```

```
console=ttyS0,115200n8"
update-grub
```

Building Container Images (takes some time ~40 minutes)

```
kolla-build --base-image 192.168.154.2:4000/ubuntu --base ubuntu
--base-tag ppc64 --type source --registry 192.168.154.2:4000
--push
```

Note: the 'gnocchi' package does not build (yet) on Power.

Note: from time to time, when the network is not reliable, some packages fail to build. In this case just restart the build process.

After the build completes you can list the images:

```
docker images
```

Cinder

The cinder implementation defaults to using LVM storage. The default implementation requires a volume group be set up. This can either be a real physical volume or a loopback mounted file for development.

During development, it may be desirable to use file backed block storage. It is possible to use a file and mount it as a block device via the loopback system.

```
mknod /dev/loop2 b 7 2
dd if=/dev/zero of=/var/lib/cinder_data.img bs=1G count=200
losetup /dev/loop2 /var/lib/cinder_data.img
pvcreate /dev/loop2
vgcreate cinder-volumes /dev/loop2
```

Deployment

Create a pair of virtual interfaces

```
ip link add type veth
brctl addif br1 veth0
ifconfig veth0 up
ifconfig veth1 up
```

Deploying Kolla (all-in-one)

All variables for the environment can be specified in the files: “/etc/kolla/globals.yml” and “/etc/kolla/passwords.yml”

Generate new passwords

```
kolla-genpwd
```

Start by editing the `globals.yml` file

```
sudo vi /etc/kolla/globals.yml

config_strategy: "COPY_ALWAYS"
kolla_base_distro: "ubuntu"
kolla_install_type: "source"
enable_haproxy: "no"
kolla_internal_vip_address: "192.168.154.2"
kolla_internal_fqdn: "{{ kolla_internal_vip_address }}"
kolla_external_vip_address: "{{ kolla_internal_vip_address }}"
kolla_external_fqdn: "{{ kolla_external_vip_address }}"
docker_registry: "192.168.154.2:4000"
network_interface: "enPlp9s0"
neutron_external_interface: "veth1"
neutron_plugin_agent: "linuxbridge"
enable_heat: "yes"
enable_neutron: "yes"
enable_cinder: "yes"
# enable_iscsi: "yes"
cinder_iscsi_ip_address: "192.168.154.2"
cinder_volume_group: "cinder-volumes"
cinder_volume_backend_name: "lvm"
```

Run the deployment

```
kolla-ansible deploy
```

After successful deployment of OpenStack, run the following command can create an `openrc` file `/etc/kolla/admin-openrc.sh` on the deploy node.

```
kolla-ansible post-deploy
```

Update configuration script

```
vi /etc/kolla/init-runonce
```

```
# Test to ensure configure script is run only once
if glance image-list | grep -q ubuntu; then
    echo "This tool should only be run once per
deployment."
    exit
fi
```



```

echo Downloading glance image.
IMAGE_URL=http://cloud-images.ubuntu.com/xenial/current/
IMAGE=xenial-server-cloudimg-ppc64el-disk1.img
if ! [ -f "$IMAGE" ]; then
    curl -L -o ./ $IMAGE $IMAGE_URL/$IMAGE
fi
echo Creating glance image.
glance image-create --name ubuntu-16.04 --progress
--disk-format qcow2 --container-format bare --progress
--file ./ $IMAGE

echo Configuring neutron.
neutron net-create public1 --router:external
--provider:physical_network physnet1
--provider:network_type flat
neutron subnet-create --name 1-subnet --disable-dhcp
--allocation-pool start=172.31.250.150,end=172.31.250.199
public1 172.31.0.0/16 --gateway 172.31.255.254
--dns-nameservers list=true 8.8.8.8
neutron net-create demo-net --provider:network_type vxlan
neutron subnet-create demo-net 10.0.0.0/24 --name
demo-subnet --gateway 10.0.0.1 --dns-nameservers list=true
8.8.8.8
neutron router-create demo-router
neutron router-interface-add demo-router demo-subnet
neutron router-gateway-set demo-router public1

```

Run the configuration script

```

source /etc/kolla/admin-openrc.sh
/etc/kolla/init-runonce

```

Logon to the Horizon dashboard

From the local system:

```
ssh -L 8080:192.168.154.2:80 root@172.31.250.2
```

Point a Web Browser at

```
http://localhost:8080/auth/login
```

```
domain: default
```

```
user: admin
```

```
password: <from /etc/kolla/admin-openrc.sh>
```

Add compute node

Logon to the system that needs to be added to the infrastructure as compute node (IP: 192.168.154.4).

```
vi /etc/rc.local
    ppc64_cpu --smt=off
```

Edit /etc/sysctl.conf

```
vi /etc/sysctl.conf
    net.ipv4.conf.all.rp_filter=0
    net.ipv4.conf.default.rp_filter=0

sysctl -p
```

Add Newton ppa (or install a different ppa if Mitaka or Kilo)

```
apt install -y software-properties-common
add-apt-repository ppa:ubuntu-cloud-archive/newton-staging
apt-get update
```

Install nova and neutron

```
apt-get install -y nova-compute
apt-get install -y neutron-linuxbridge-agent
```

Stop the services

```
service nova-compute stop
service neutron-linuxbridge-agent stop
```

Configure the services

Download some configuration files from the Kolla controller system

```
cd ~
scp root@192.168.154.2:/etc/kolla//nova-compute/nova.conf .
scp
root@192.168.154.2:/etc/kolla//neutron-linuxbridge-agent/neutron
.conf .
scp
root@192.168.154.2:/etc/kolla//neutron-linuxbridge-agent/ml2_con
f.ini .
```

Update the configuration files

```
vi ml2_conf.ini

[ml2]
```

```

type_drivers = flat,vlan,vxlan
tenant_network_types = vxlan
mechanism_drivers = linuxbridge,l2population

[ml2_type_vlan]
network_vlan_ranges =

[ml2_type_flat]
flat_networks = physnet1

[ml2_type_vxlan]
vni_ranges = 1:1000
vxlan_group = 239.1.1.1

[securitygroup]
firewall_driver =
neutron.agent.linux.iptables_firewall.IptablesFirewallDriver

[linux_bridge]
physical_interface_mappings = physnet1:veth1

[vxlan]
l2_population = true
local_ip = 192.168.154.4

```

vi neutron.conf

```

[DEFAULT]
debug = False
log_dir = /var/log/kolla/neutron
use_stderr = False
# bind_host = 192.168.154.2
bind_host = 0.0.0.0
bind_port = 9696
api_paste_config = /usr/share/neutron/api-paste.ini
endpoint_type = internalURL
metadata_proxy_socket =
/var/lib/neutron/kolla/metadata_proxy
interface_driver =
neutron.agent.linux.interface.BridgeInterfaceDriver
allow_overlapping_ips = true
core_plugin = ml2
service_plugins = router

[nova]
auth_url = http://192.168.154.2:35357
auth_type = password
project_domain_id = default

```

```

user_domain_id = default
region_name = RegionOne
project_name = service
username = nova
password = KXZU00G0p68HukxqHq8UI86U7kxzJuSq8HGvbixK
endpoint_type = internal

[oslo_concurrency]
lock_path = /var/lib/neutron/tmp

[oslo_messaging_rabbit]
rabbit_userid = openstack
rabbit_password = bstpqHG52O8sfYl21JYWfz7ReA8Jg6KXmr9bggcA
rabbit_ha_queues = true
rabbit_hosts = 192.168.154.2:5672

[agent]
root_helper = sudo neutron-rootwrap
/etc/neutron/rootwrap.conf

[database]
connection =
mysql+pymysql://neutron:Xl07xeMLcsmubQma277RfOeh8ULRJYomn9w
698oj@192.168.154.2:3306/neutron
max_retries = -1

[keystone_auth_token]
auth_uri = http://192.168.154.2:5000
auth_url = http://192.168.154.2:35357
auth_type = password
project_domain_id = default
user_domain_id = default
project_name = service
username = neutron
password = TkK0fbFRm4IFFMu1Nt93r224o1XhqPej1CE9f5HY
memcache_security_strategy = ENCRYPT
memcache_secret_key =
mojgGPLte3GcJ5dto67zXhrEWlVOpreBRdf1Ev4l
memcached_servers = 192.168.154.2:11211

[oslo_messaging_notifications]
driver = noop

```

```

vi nova.conf
[DEFAULT]
debug = False
log_dir = /var/log/kolla/nova
use_forwarded_for = true
api_paste_config = /etc/nova/api-paste.ini

```

```
state_path = /var/lib/nova
osapi_compute_listen = 0.0.0.0
osapi_compute_listen_port = 8774
osapi_compute_workers = 4
metadata_listen = 0.0.0.0
metadata_listen_port = 8775
metadata_workers = 4
ec2_listen = 0.0.0.0
ec2_listen_port = 8773
use_neutron = True
firewall_driver = nova.virt.firewall.NoopFirewallDriver
scheduler_max_attempts = 10
linuxnet_interface_driver =
nova.network.linux_net.NeutronLinuxBridgeInterfaceDriver
allow_resize_to_same_host = true
compute_driver = libvirt.LibvirtDriver
my_ip = 192.168.154.4

[vnc]
novncproxy_host = 0.0.0.0
enabled = True
novncproxy_port = 6080
vncserver_listen = 0.0.0.0
vncserver_proxyclient_address = 192.168.154.4
novncproxy_base_url =
http://192.168.154.2:6080/vnc_auto.html

[oslo_messaging_rabbit]
rabbit_userid = openstack
rabbit_password = bstpqHG52O8sfYl21JYWfz7ReA8Jg6KXmr9bggcA
rabbit_ha_queues = true
rabbit_hosts = 192.168.154.2:5672

[oslo_concurrency]
lock_path = /var/lib/nova/tmp

[glance]
api_servers = http://192.168.154.2:9292
num_retries = 1

[cinder]
catalog_info = volume:cinder:internalURL

[neutron]
url = http://192.168.154.2:9696
auth_strategy = keystone
metadata_proxy_shared_secret =
pBsCzjrZ8XGTEWLurBm8ywsEPw8T130RusiLOR2q
service_metadata_proxy = true
```

```
auth_url = http://192.168.154.2:35357
auth_type = password
project_domain_name = default
user_domain_id = default
project_name = service
username = neutron
password = TkK0fbFRm4IFFMulNt93r224o1XhqPej1CE9f5HY

[database]
connection =
mysql+pymysql://nova:5GnBh03QoXQP3Ly1mL30UH27x7GBhRFJ2akyeq
Km@192.168.154.2:3306/nova
max_pool_size = 50
max_overflow = 1000
max_retries = -1

[api_database]
connection =
mysql+pymysql://nova_api:4sfRFI8uOH094YDmhVj8Wgoo5sWAIkFRqY
hrshta@192.168.154.2:3306/nova_api
max_retries = -1

[cache]
backend = oslo_cache.memcache_pool
enabled = True
memcache_servers = 192.168.154.2:11211

[keystone_authtoken]
auth_uri = http://192.168.154.2:5000
auth_url = http://192.168.154.2:35357
auth_type = password
project_domain_id = default
user_domain_id = default
project_name = service
username = nova
password = KXZU00G0p68HukxqHq8UI86U7kxzJuSq8HGvbixK
memcache_security_strategy = ENCRYPT
memcache_secret_key =
mojgGPLte3GcJ5dto67zXhrEWlVOpReBRdf1Ev4l
memcached_servers = 192.168.154.2:11211

# [libvirt]
# connection_uri = "qemu+tcp://192.168.154.4/system"

[upgrade_levels]
compute = auto

[oslo_messaging_notifications]
driver = noop
```

```
[conductor]
workers = 4

vi nova-compute.conf

[DEFAULT]
[libvirt]
virt_type=qemu
```

Load the configuration files

```
cat ml2_conf.ini >>
/etc/neutron/plugins/ml2/linuxbridge_agent.ini

cp /etc/neutron/neutron.conf /etc/neutron/neutron.conf.orig

cp neutron.conf /etc/neutron/neutron.conf

cp /etc/nova/nova.conf /etc/nova/nova.conf.orig

cp nova.conf /etc/nova/nova.conf

cp nova-compute.conf /etc/nova/nova-compute.conf
```

Start the service manually

```
/usr/bin/python /usr/bin/nova-compute
--config-file=/etc/nova/nova.conf
--config-file=/etc/nova/nova-compute.conf
--log-file=/var/log/nova/nova-compute.log

/usr/bin/python /usr/bin/neutron-linuxbridge-agent
--config-file=/etc/neutron/neutron.conf
--config-file=/etc/neutron/plugins/ml2/linuxbridge_agent.ini
--log-file=/var/log/neutron/neutron-linuxbridge-agent.log
```

Kill the services

Restart the services

```
service nova-compute start

service neutron-linuxbridge-agent start
```

Troubleshooting

If nova-compute logs (docker logs nova-compute) show a 'permission denied error' on KVM:

```
chmod 777 /dev/kvm
```

Test the registry

```
wget http://192.168.154.2:4000/v2/_catalog
wget http://192.168.154.2:4000/v2/ubuntu/tags/list

docker pull 192.168.154.2:4000/ubuntu:ppc64
ppc64: Pulling from ubuntu
Digest:
sha256:2af24e0b4f901bfe26ca059c28099b939287c7c73bbe8c26b1dfc06b3
472f2ba
Status: Image is up to date for 192.168.154.2:4000/ubuntu:ppc64

docker run -t -i 192.168.154.2:4000/ubuntu:ppc64 bash
root@e142a20f5597:/#
```

Change global IP addresses

```
find ./ -type f -exec sed -i 's/192\.168\.154\.2/10\.154\.10\.212/g'
{} \;
```

Test Cinder

```
openstack volume create --size 1 frb-test-vol
```

Field	Value
attachments	[]
availability_zone	nova
bootable	false
consistencygroup_id	None
created_at	2016-11-18T14:50:54.735141
description	None
encrypted	False
id	2a442bcc-533c-4909-9912-283acc040c5b
migration_status	None
multiattach	False
name	frb-test-vol
properties	
replication_status	disabled
size	1
snapshot_id	None
source_volid	None
status	creating
type	None
updated_at	None
user_id	24bf3349e7bb4cefbdd6c8a407003f76

+-----+-----+

openstack volume list

ID	Display Name	Status	Size	Attached to	
2a442bcc-533c-4909-9912-283acc040c5b	frb-test-vol	available	1		